

Real-time Hand Tracking with Neural Nets on an FPGA

Brian Kubisiak <bkubisia@caltech.edu>
Quinn Osha <qosha@caltech.edu>

1 Functional Specification

For our project, we will be designing and implementing a system for tracking a hand in real time from a camera input. The system will be implemented using a Digilent Nexys 4 FPGA development board. It will take an input from a camera connected to the USB port, analyze the data on the Artix 7 FPGA, and output each frame over the VGA port with a cursor overlay. The cursor overlay will indicate where in the frame the hand is located. While the neural net is computing the hand position, the frame will temporarily be stored in the board's memory—the FPGA is too small to process an entire frame at once.

1.1 Operation

The FPGA design will consist of three distinct layers:

Input Layer For communicating with the USB camera, writing the frame to memory, and converting the frame into a format that the neural net can use.

Neural Network For analyzing the camera data to determine where the hand is in the frame.

Output Layer For reading the frame back from memory, overlaying a cursor on the hand position, and generating control signals to output the frame through a VGA port.

A top-level block diagram showing the layers and their connections is shown in figure (1). The Rx and Tx inputs come from the USB-UART bridge. The Data, Address, Wr, and Rd signals go to the memory on the Digilent board. Note that the Data and Address buses must be muxed before connecting to the pins. The HSYNC, VSYNC, and Data outputs will go to the VGA port on the board.




Figure 1: Top-level block diagram for the system.

1.1.1 Input Layer

1.1.2 Neural Network

1.1.3 Output Layer

The USB controller will load each pixel serially from the camera, storing the pixel in a shift register. Once the pixel is completely shifted into the controller, the controller will output the pixel (in parallel) to the input frame buffer.

This input frame buffer will act like a FIFO for loading data from the USB controller. Once the buffer is full, it will output the data in parallel to both the neural net and the output buffer. Once this data is sent to the neural net and output buffer for processing, the USB controller can begin loading new data into the input buffer.

The output buffer will have two inputs: the first input will be a massively parallel input for loading all the data from the input buffer. The second input will allow individual pixels to be reset. This will be connected to the cursor drawing logic in order to draw the cursor at the hand's location.

The neural network will load all the data in parallel from the input buffer and operate on the data using the machine-learned parameters to find the coordinates of the hand that is being tracked. These coordinates will then be output to the cursor drawing logic. The cursor drawing logic will accept the coordinates produced by the neural net and will output control signals to the output buffer for drawing the cursor.

The output buffer—after loading the input data and having the cursor drawn to it—will output one pixel at a time to the VGA controller. The VGA controller will load this data into a shift register, then shift it out to the display (along with the needed control signals). Due to the separate input and output buffers, a frame can be input while the previous frame is being shifted out.

1.2 Algorithms

We will be using a convolutional neural network to identify where the hand is located in each frame. The parameters for the neural network will be learned in advance on the computer, then hard-coded into the FPGA neural net. Time-permitting, we will also implement unsupervised learning on the FPGA in order to allow the system to learn without hard-coded parameters.

1.3 Inputs

The input to this system will be video input through a USB controller. Each frame that is input to the USB controller will be stored in an internal frame buffer before being processed by the neural net.

1.4 Outputs

The output from this system will be through a VGA controller. The VGA output will be taken from an output frame buffer. The output frame will be generated from the input frame with a cursor overlaid on the position of the hand (identified by the neural net).

2 Schedule

Date	Tasks
4/7–4/13	Research current implementations of neural nets on FPGAs. Decide on video input and output devices (size and pixel depth). Write code for USB controller.
4/14–4/20	Write code for input data buffer and input data handling. Design—in a high-level-language—the desired neural net.
4/21–4/27	Test and configure the neural net on the computer. Write code for output data buffer and handling.
4/28–5/4	Begin coding the neural net on the FPGA. Write code for output display handling using VGA.
5/5–5/11	Finish coding the neural net.
5/12–5/18	Write code for the cursor video output. That is, given x and y coordinates, print the cursor at that point.
5/19–5/25	Write the top-level entity for the system, combining all the blocks. Train (unsupervised) the neural net.
5/26–6/1	Finish training the neural net. Fix any other bugs/issues.
6/2–6/5	Finish documentation and final touches.

3 Demonstration

This system will be demonstrated by attaching a video camera to the USB input and a monitor to the VGA output. Then, the system should detect a hand in the input frame and output the frame to the monitor with a cursor overlaid on the position of the hand.